QuantRule - Reproducible Trading Evidence

QuantRule

2026-11-12

Users

- **One quantitative researcher** - the only human user. Works alone, on one machine or a small internal server.
- **One reviewer per published result** - not a system user. Receives an evidence bundle and a digest, and reproduces the result offline, without an account.

Pattern: short interactive sessions, between long unattended runs that continue after the client closes.

Requirements

- **R1.** Ingest local M1 OHLCV candles; identify each dataset by a content hash.
- **R2.** Build a multi-timeframe ladder; publish readings, states, events, and progression as numbers, never labels.
- **R3.** Evaluate one human-authored rule through strategy, execution, and portfolio, with explicit costs and no look-ahead.
- **R4.** Run a backtest as a job that survives the client closing and is readable from another machine.
- **R5.** Produce a reproducible bundle with provenance and a comparable digest.
- **R6.** Expose the same capabilities through a browser client and a command shell.
- **R7.** Support chart-first exploration: a human marks good and bad entries on a window, a deterministic oracle marks the same window, and the check runs on a later, unmarked window.
- **R8.** Keep candle data on the operator's machine and require no account.

Additional Context

- Single operator, small deployment: one machine or one internal server.
- The goal is a durable, cost-aware edge on OHLCV candles. "No edge" is a valid result.
- Determinism is the trust claim: the same inputs produce the same bytes.
- Runs are batch and CPU-bound: minutes to hours, not milliseconds.