QuantRule - Finding a Trading Edge

QuantRule

2026-11-12

Users

One quantitative researcher - the only user. A systematic trader working alone, on their own OHLCV candles, on one machine or a small internal server.

- **What they are looking for:** an *edge* - a measurable advantage that makes a strategy produce positive expected value over many trades. A 45% win rate with +3R winners and -1R losers is an edge: `(0.45 x 3) - (0.55 x 1) = +0.8R` per trade before costs. A high win rate with a poor payoff is not.
- **How they work:** short interactive sessions - marking a chart, inspecting a run - between long unattended runs that continue after the client closes. Many hypotheses, few survivors.
- **What they need:** to kill a hypothesis cheaply and defend a survivor months later; a negative result must be as trustworthy as a positive one.
- **Why one user:** accounts, sharing, and tenancy would add cost without improving the research.

Reader of the evidence - not a user. A reviewer receives a bundle and a digest, and reproduces the run offline, without an account.

Pattern: many hypotheses, few survivors.

Requirements

- **R1.** Ingest local M1 OHLCV candles; identify each dataset by a content hash.
- **R2.** Build a multi-timeframe ladder; publish readings, states, events, and progression as numbers, never labels.
- **R3.** Evaluate one human-authored rule through strategy, execution, and portfolio, with explicit costs and no look-ahead; report expectancy and the inputs that produced it.
- **R4.** Run a backtest as a job that survives the client closing and is readable from another machine.
- **R5.** Produce a reproducible bundle with provenance and a comparable digest.
- **R6.** Expose the same capabilities through a browser client and a command shell.
- **R7.** Support chart-first exploration: a human marks good and bad entries on a window, a deterministic oracle marks the same window, and the check runs on a later, unmarked window.
- **R8.** Keep candle data on the operator's machine and require no account.

Additional Context

- Single operator, small deployment: one machine or one internal server.
- The goal is an edge. "No edge" is the most common outcome and a valid result. A backtest is a filter, not proof: out-of-sample data, costs, and a large enough sample decide.
- Determinism is the trust claim: the same inputs produce the same bytes. Runs are batch and CPU-bound: minutes to hours, not milliseconds.