Checkpoint A - Context and First Behaviour

QuantRule

2026-11-12

Kata

See `docs/course/architecture/src/01-kata.md`. Eight numbered requirements (R1-R8), one operator, two clients, no accounts.

Claim

Claim type: Behaviour.

> For a quantitative researcher working on their own OHLCV data, the system will produce one evidence bundle per submitted rule, measured by comparing the canonical digest of two runs over the same inputs, with the limitation that a digest says nothing about whether the rule has an edge.

Falsified by: two runs over identical inputs and the same image producing different artifact bytes.

First increment

One user action, one observable result:

```
backtest run --config <run> --strategy <rule> --out <bundle>
  -> exit 0, bundle written, provenance.json names data version and config hash
```

It is implemented today as `backtest run` in the v0.3.0 tree and is being moved behind the distributed run lifecycle. It already crosses the architecture boundary: shell -> application service -> engines, with the engines pure.

Evidence it carries: a printed-output shell test, an integration test over fixture candles, and the run manifest.

Pitch (45-60 s)

- **QuantRule** - a research workstation that turns a candle hypothesis into evidence you can reproduce.
- **The user** is one researcher who must defend a result months later.
- **The claim**: one rule in, one bundle out, byte-identical on a rerun.

Panel

- **Name:** Quantitative researcher.
- **Project:** QuantRule - reproducible OHLCV evidence.
- **Description (<= 10 words):** Reproducible evidence for rule-based OHLCV research.

- **Claim:** Same inputs, same bundle, every time.