Checkpoint A - Context and First Behaviour

2026-11-12

Context

One quantitative researcher works alone on their own OHLCV data. The system must ingest candles, measure them across timeframes as numbers, evaluate one human-authored rule without look-ahead, and produce evidence a reviewer can reproduce months later. Runs take minutes to hours and must survive the client closing. Candle data stays on the operator's machine; there are no accounts.

Claim

Claim type: Behaviour.

> For a quantitative researcher working on their own OHLCV data, the system will produce one evidence bundle per submitted rule, measured by comparing the canonical digest of two runs over the same inputs, with the limitation that a digest says nothing about whether the rule has an edge.

Falsified by: two runs over identical inputs and the same image producing different artifact bytes.

First increment

One user action, one observable result:

```
submit a rule over a dataset
  -> a bundle is written
  -> provenance records the data version, the config hash, the model version,
     and the image digest
  -> the command exits 0
```

It crosses the real architecture boundary: client -> service -> engines, with the engines pure and the evidence written by one owner.

Evidence: a printed-output client test, an integration test over fixture candles, and the bundle's provenance file. Limit: fixture data only; no out-of-sample claim.

Pitch (45-60 seconds)

- **QuantRule** - a research workstation that turns a candle hypothesis into evidence you can reproduce.
- **The user** is one researcher who must defend a result months later.
- **The claim**: one rule in, one bundle out, byte-identical on a rerun.

Panel

- **Name:** Quantitative researcher.
- **Project:** QuantRule - reproducible OHLCV evidence.
- **Description (10 words):** Reproducible evidence for rule-based OHLCV research.
- **Claim:** Same inputs, same bundle, every time.