Checkpoint B - Operating Model and Feedback

QuantRule

2026-11-12

Lifecycle

Iterative and incremental, with pre-registered increments.

- Dominant uncertainty: whether a rule-based edge exists at all. No lifecycle choice removes that uncertainty; only disciplined experiments measure it.
- Therefore each increment is a frozen, pre-registered run: hypothesis, data version, costs, and green criteria recorded before execution.
- Iteration is allowed on the platform; it is not allowed on a hypothesis after it has been run.

Evidence: `docs/workflows/improvement-workflow.md` and the run manifest contract. Limit: the model has no accepted baseline yet, so the loop has not yet closed once.

Flow of one item

```
hypothesis -> pre-registration -> data version -> run -> review -> decision
    ↑                                                              |
    +------------ new hypothesis (new record) ---------+
```

Stage	Waits	Depends on
Hypothesis	on the human	one written sentence and a rule
Pre-registration	on review	data version, costs, green criteria
Run	on the machine	frozen config; minutes to hours
Review	on the human	the bundle and its digest
Decision	on the human	review, not on the run

The longest wait is a run, and it is unattended by design.

Feedback

- Who: the human, reading the review and the digest.
- Which decision it can change: accept, revise, or stop the hypothesis.
- Who can revise code before commitment: the human, before the config is frozen; after the run, a change is a new pre-registration.

Unit and done rule

- **Unit:** one issue with one observable acceptance criterion.
- **Done:** the gate passes, the change is merged, and the evidence is recorded.
- **Stop rule:** a unit that grows past one deliverable is split, not extended.

Review rule

Peer-style review of the diff plus the gate (`cargo fmt`, `clippy -D warnings`, `cargo test`, documentation checkers). A red gate blocks the merge; a review comment blocks it too.

MVP or incremental

Incremental. The value is established - a reproducible backtest is useful on its own - and the scope grows by adding services, not by testing whether users want the product.

What would revise the plan

A run that cannot be reproduced from its own provenance, or a platform change that forces a hypothesis to be re-run under new costs. Either would stop feature work and fix the evidence chain first.

Observed: the pipeline and the gate exist. **Proposed:** the distributed lifecycle. **Unknown:** whether any rule survives out-of-sample.