Checkpoint B - Operating Model and Feedback

2026-11-12

Lifecycle

Iterative and incremental, with pre-registered increments.

- Dominant uncertainty: whether a rule-based edge exists at all. No lifecycle removes that; only disciplined experiments measure it.
- Each increment is a frozen run: hypothesis, data version, costs, and green criteria recorded before execution.
- The platform may iterate freely. A hypothesis may not: after it runs, a change is a new, separately recorded hypothesis.

Observed: the pipeline, the bundle contract, and the gate exist. **Proposed:** the distributed run lifecycle. **Unknown:** whether any rule survives out-of-sample.

Flow of one item

```
hypothesis -> pre-registration -> data version -> run -> review -> decision
     ^                                                             |
     +------------ new hypothesis (new record) --------------+
```

Stage	Waits on	Depends on
Hypothesis	the human	one written sentence and one rule
Pre-registration	review	data version, costs, green criteria
Run	the machine	a frozen config; minutes to hours
Review	the human	the bundle and its digest
Decision	the human	the review, not the run

The longest wait is a run, and it is unattended by design.

Feedback

- Who supplies it: the human, reading the review and the digest.
- Which decision it can change: accept, revise, or stop.
- Who can revise code before commitment: the human, before the config is frozen. After the run, a change is a new pre-registration.

Unit and done rule

- **Unit:** one issue with one observable acceptance criterion.
- **Done:** the gate passes, the change is merged, and the evidence is recorded.
- **Stop rule:** a unit that grows past one deliverable is split, not extended.

Review rule

Defined before work starts: review of the diff plus the gate (`fmt`, `clippy -D warnings`, tests, documentation checks). A red gate blocks the merge; a review comment blocks it too.

MVP or incremental

Incremental. The value is established - a reproducible backtest is useful on its own - so scope grows by adding services, not by testing whether anyone wants the product.

What would revise the plan

A run that cannot be reproduced from its own provenance, or a platform change that forces a completed hypothesis to be re-run under new costs. Either stops feature work until the evidence chain is fixed.